

Where does my data actually live?

Three homes for your data, and how they pass it back and forth.

Excel is no longer the end of the road. It's one stop on a round trip through the cloud. Today's three "homes" each play a different role: **Excel** is where you do the work, **Power BI** is where the trusted numbers live, and **Dataverse** is where the full business records live. Knowing which home you're pulling from explains everything you'll see in the demo.

X**THE WORKSHOP****Excel workbook**

Where you build the data model, write measures, shape it, and explore. Lives on your machine or OneDrive.

Gives you: full control, fast iteration, but no shared truth.



PUBLISH
OR PULL

B**THE TRUSTED VIEW****Power BI semantic model**

A curated dataset in the cloud.
Relationships, measures, and business logic baked in. One version of the truth.

Gives you: clean numbers, refreshable, governed.



CONNECT
AS A
TABLE

D**THE SYSTEM OF RECORD****Dataverse table**

An enterprise object. Tracks ownership, permissions, relationships, workflow status. The full record, not just numbers.

Gives you: complete business context, system columns, GUIDs.

The round trip we'll demo today**1**

Build a data model in Excel (relationships and measures).

2

Publish that model up to Power BI as a semantic model.

3

Pull the governed semantic model back into Excel.

4

Pull a Dataverse table into Excel and see what enterprise data looks like.

Three ways to get data into Excel, compared

Same destination (Excel), very different starting points. Here's what changes.

THE CLASSIC

From a CSV file

WHAT YOU'RE GETTING

A raw text file. Just rows and columns.

WHO OWNS THE LOGIC

You do. Cleaning, types, relationships, measures: all on you.

WHEN DATA CHANGES

You re-import the file and redo your shaping.

WHAT IT LOOKS LIKE

Plain columns. No system metadata. No relationships.

USE IT WHEN

You're doing one-off analysis or there is no governed source.

THE CURATED VIEW

From a Power BI semantic model

WHAT YOU'RE GETTING

A live connection to a published, certified dataset.

WHO OWNS THE LOGIC

The model does. Relationships and measures are already baked in.

WHEN DATA CHANGES

Hit Refresh. Excel just asks Power BI for the latest numbers.

WHAT IT LOOKS LIKE

Clean field names. Pre-built measures. Shared definitions everyone agrees on.

USE IT WHEN

You want trusted numbers without rebuilding the logic.

THE SYSTEM OF RECORD

From a Dataverse table

WHAT YOU'RE GETTING

The whole enterprise object, including system columns.

WHO OWNS THE LOGIC

The platform. It tracks owners, permissions, status, GUIDs.

WHEN DATA CHANGES

Refresh. The connection stays live to the table.

WHAT IT LOOKS LIKE

Lots of fields. `_createdby_value`, `_ownerid_value`, GUIDs. That's normal. Trim with Power Query.

USE IT WHEN

You need real business records with full context, not just analytics.

The mental model: CSV gives you a snapshot you have to babysit. Power BI gives you a clean, curated view someone has already polished. Dataverse gives you the whole record, system columns and all. Same Excel destination, very different starting points.

So what does this have to do with AI?

Day 1 is the foundation. The rest of the week sits on top of it.

Where Day 1 fits in the modern Excel AI stack

Day 5 · Copilot Studio
Custom copilots and governance

Day 4 · Office Scripts & Power Automate
Automation and workflows

Day 3 · Excel: Copilot, Python, Modern formulas
Analyze, visualize, act on data

Day 2 · Power Query & Dataflows
Clean and prepare data for analysis

Day 1 · Power BI & Dataverse
Data backbone (you are here)

Every layer above this one assumes the data underneath is trustworthy. Get this layer right and Copilot, Python in Excel, and your custom agents have something solid to stand on.

Why governed data is the unlock for AI

AI is only as good as the data it sees

Copilot and Python in Excel can't fix mystery columns, broken joins, or "which revenue number is right." Garbage in, confidently-wrong out. The data layer determines the ceiling on AI quality.

Semantic models teach AI the meaning of your data

A semantic model isn't just numbers. It's labels, relationships, and definitions. That context is exactly what an LLM needs to answer "what was Q3 margin in West?" without guessing.

Dataverse gives AI a consistent system of record

When agents act on your behalf (Day 5), they need real records to read and update, not loose CSVs. Dataverse provides that backbone.

Refreshable beats re-imported

A live connection means today's prompt works tomorrow on tomorrow's data. No re-pasting. No forgotten cleaning steps.

Bottom line: Better data foundations mean better prompts, better answers, and far fewer "the AI made it up" moments. Day 1 isn't a detour from AI. It's the part of AI most people skip, then wonder why their copilots are flaky.