

From a messy file to reusable cloud logic

The Day 2 journey: clean once in Excel, then push that same cleanup to the cloud so everyone (and every AI tool) gets the same clean table.

Today's demo file ([quarterly_claims.xlsx](#)) is "fine for humans, terrible for AI." Quarterly columns spread across the sheet, missing customer names, agent info crammed into one cell. We'll fix it three times, on purpose. First by hand in Power Query, then with Copilot's help, and finally by promoting the same logic up to a cloud Dataflow. Each step does something the one before it can't.

1 THE MESS

Raw Excel file

Quarter columns sprawled across the sheet. Agent_Info jammed into one column. Blank customer names. Looks fine, refuses to be analyzed.

Try Analyze Data on it. Copilot has no idea what to do.

→
CLEAN
LOCALLY

2 THE CLEANUP

Power Query in Excel

Unpivot the quarter columns. Split Agent_Info. Fill down customer names. Set types. Save the M code as your repeatable recipe.

Copilot helps you plan the cleanup and polish the M code (more on page 2).

→
PROMOTE
TO
CLOUD

3 THE REUSE

Dataflow in the cloud

Lift the same M code into a Dataflow at [make.powerapps.com](#). Land the cleaned table in Dataverse. Schedule refresh.

Now Excel, Power BI, Copilot Studio, and Power Automate all read the same clean table.

Why each step exists

Step 1: feel the pain

Try to ask Copilot something useful about the raw file. It can't answer. That's the proof we need to clean.

Step 2: build the recipe

Power Query records every click as M code. That recipe is what we'll reuse, not the cleaned file itself.

Step 3: lift, don't repeat

Copy that M code into a Dataflow. Same logic, now running in the cloud on a schedule.

Step 4: feed the platform

The cleaned table sits in Dataverse, ready for any tool that needs governed data.

What is Copilot actually doing in Power Query?

Copilot is a thinking partner around Power Query, not a button inside it. Here's what that means in practice.

WHAT COPILOT CAN DO FOR YOU

Thinking partner

- + **Diagnose the data.** Paste a sample and ask "what's wrong with this table?" Copilot will spot blanks, inconsistent types, columns that should be unpivoted.
- + **Sketch a cleanup plan.** Ask for a step-by-step Power Query plan. You get an outline before you start clicking.
- + **Rewrite your M code.** Paste the script. Copilot adds comments, renames cryptic steps, and tightens the structure (without changing the logic).
- + **Explain unfamiliar M.** Inherited a query? Ask Copilot to translate it into plain English so you actually understand it.

WHAT COPILOT DOES NOT DO

Hands-on driver

- ✗ **It does not click inside Power Query.** It cannot push the Unpivot button or change a data type for you.
- ✗ **It does not always catch the real fix.** On this dataset it will suggest formatting tweaks but often miss that quarterly columns need to be unpivoted. You still have to know the move.
- ✗ **It does not validate the output.** You are still the one who checks the final table looks right.
- ✗ **It does not move logic to the cloud.** Promoting a query to a Dataflow is your job, not Copilot's.

Three concrete prompts to use during the demo

PROMPT 1: REVIEW

Find the issues

"Here is a table called `quarterly_claims`. Review it for any potential data quality issues. Don't fix anything yet. Just tell me what looks off."

PROMPT 2: PLAN

Outline the cleanup

"Based on what you noticed, outline a short step-by-step cleaning plan I could apply in Power Query."

PROMPT 3: POLISH

Tighten the M code

"Review the steps I've applied so far and suggest improvements to cleanliness, clarity, and structure. Don't change the logic."

So what does pushing this to a Dataflow actually do?

Same M code, very different reach. Here's what changes when the cleanup logic moves to the cloud.

LOCAL

Power Query inside Excel

WHERE THE LOGIC LIVES

Inside one workbook on your machine or OneDrive.

WHO CAN USE IT

You, and anyone you send the file to.

WHEN THE SOURCE UPDATES

You open the file and click Refresh.

WHERE THE CLEAN TABLE LANDS

A sheet in your workbook.

WHAT OTHER TOOLS SEE

Nothing. The cleanup is invisible outside this file.

BEST FOR

One-off analysis, exploration, your own reporting.

CLOUD

The same logic as a Dataflow

WHERE THE LOGIC LIVES

In Power Query Online at make.powerapps.com. One central definition.

WHO CAN USE IT

Anyone in your tenant with permission. Excel, Power BI, Copilot Studio, Power Automate, Dataverse.

WHEN THE SOURCE UPDATES

Scheduled refresh runs in the cloud automatically. No "send me the latest file."

WHERE THE CLEAN TABLE LANDS

A Dataverse table that any platform tool can query.

WHAT OTHER TOOLS SEE

A governed, consistent table they can all build on.

BEST FOR

Anything that needs to be reused, refreshed, or fed into automation and AI.

The mental model: Power Query in Excel cleans data for *this one file*. A Dataflow takes that exact same cleanup and runs it for *everyone*, on a schedule, in the cloud. You're not redoing the work. You're promoting it.

Why this matters for AI: Copilot Studio agents and Power Automate flows need a stable, named table to read from. A Dataflow gives them exactly that. Skip this layer and every AI tool downstream is guessing at messy inputs. Build it once and the rest of the week (Copilot in Excel, Office Scripts, Copilot Studio) all stand on solid ground.